

inst.eecs.berkeley.edu/~cs61c

CS61C : Machine Structures

Lecture 11 – RISC-V Instruction Representation II

Instructors
Dan Garcia and Bora Nikolic
2018-09-17

Developers, Despair: Half Your Time Is Wasted on Bad Code

A study by online payments firm Stripe found 31% of developer time is wasted on routine maintenance work, draining \$300 billion each year from global GDP. The survey of 1,000 developers and 1,000 C-level executives found that on average, about half of a developer's work time is devoted to maintenance activities such as debugging, modifying, and fixing code. Developers reported working 41 hours per week on average, including 17.3 hours on maintenance work.

Liam Tung, ZDNet, Sept 11, 2018.



Review

- **Simplification works for RISC-V: Instructions are same size as data word (one word) so that they can use the same memory.**
- **Computer actually stores programs as a series of these 32-bit numbers.**
- **RISC-V Machine Language Instruction:**
32 bits representing a single instruction
 - We covered R-type, I-type and S-type instructions
 - Learned a few more new instructions
 - `slt rd, rs1, rs2` - set less than: $R[rd] = (R[rs1] < R[rs2]) ? 1 : 0$
 - `sltu, slti, sltiu`
 - `lh, lhu, sh`



Levels of Representation/Interpretation

High Level Language
Program (e.g., C)

Compiler

Assembly Language
Program (e.g., RISC-V)

Assembler

Machine Language
Program (RISC-V)

*Machine
Interpretation*

Hardware Architecture Description
(e.g., block diagrams)


*Architecture
Implementation*

Logic Circuit Description
(Circuit Schematic Diagrams)


```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

```
lw x10, 0(x12)  
lw x11, 4(x12)  
sw x11, 0(x12)  
sw x10, 4(x12)
```

Anything can be represented
as a *number*,
i.e., data or instructions



1100	0110	1010	1111	0101	1000
0101	1000	0000	1001	1100	0110
1010	1111	0101	1000	0000	1001



RISC-V Conditional Branches

- E.g., **BEQ x1, x2, Label**
- **Branches read two registers but don't write a register (similar to stores)**
- **How to encode label, i.e., where to branch to?**



Branching Instruction Usage

- **Branches typically used for loops (`if-else`, `while`, `for`)**
 - Loops are generally small (< 50 instructions)
 - Function calls and unconditional jumps handled with jump instructions (J-Format)
- **Recall: Instructions stored in a localized area of memory (Code/Text)**
 - Largest branch distance limited by size of code
 - Address of current instruction stored in the program counter (PC)



PC-Relative Addressing

- **PC-Relative Addressing:** Use the `immediate` field as a two's-complement offset to PC
 - Branches generally change the PC by a small amount
 - Can specify $\pm 2^{11}$ 'unit' addresses from the PC
 - (We will see in a bit that we can encode 12-bit offsets as immediates)
- **Why not use byte as a unit of offset from PC?**
 - Because instructions are 32-bits (4-bytes)
 - We don't branch into middle of instruction



Scaling Branch Offset

- **One idea: To improve the reach of a single branch instruction, multiply the offset by four bytes before adding to PC**
- **This would allow one branch instruction to reach $\pm 2^{11} \times 32\text{-bit}$ instructions either side of PC**
 - **Four times greater reach than using byte offset**



Branch Calculation

- If we **don't** take the branch:

$PC = PC + 4$ (i.e., next instruction)

- If we **do** take the branch:

$PC = PC + \text{immediate} * 4$

- Observations:

- **immediate** is number of instructions to jump (remember, specifies words) either forward (+) or backwards (−)

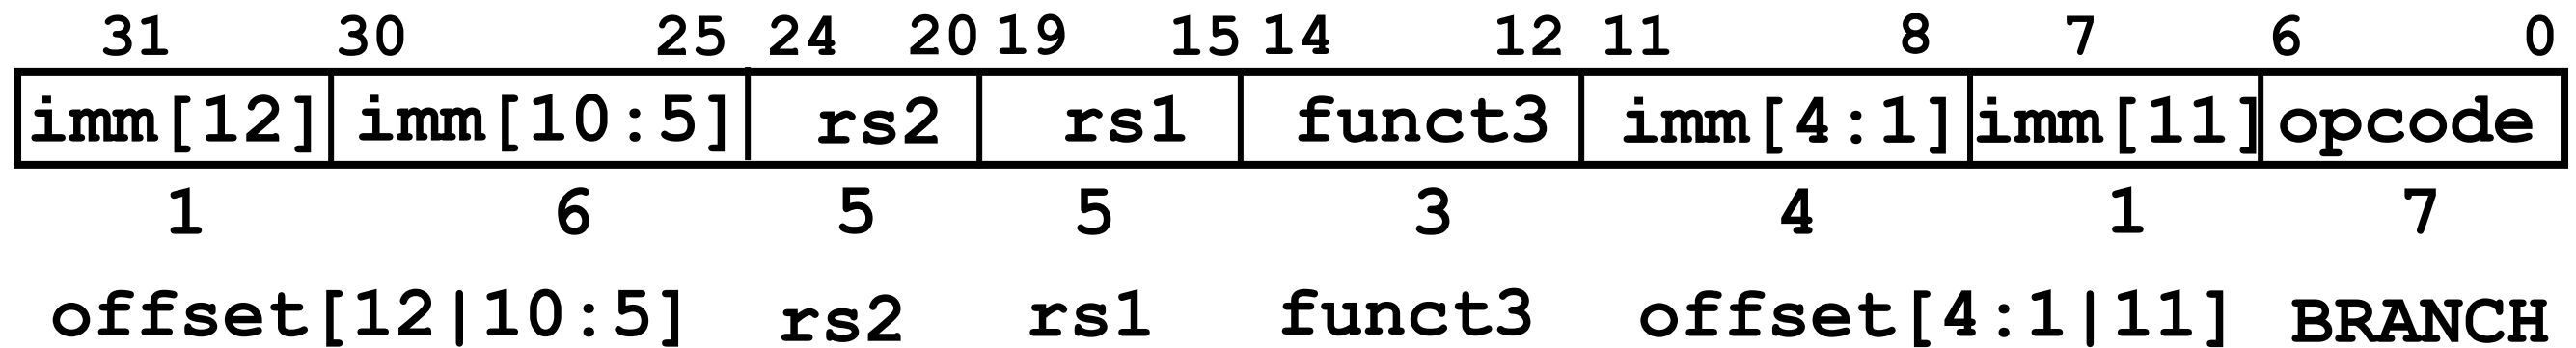


RISC-V Feature, $n \times 16$ -bit instructions

- Extensions to RISC-V base ISA support 16-bit compressed instructions and also variable-length instructions that are multiples of 16-bits in length
- To enable this, RISC-V scales the branch offset by 2 bytes even when there are no 16-bit instructions
- Reduces branch reach by half and means that $\frac{1}{2}$ of possible targets will be errors on RISC-V processors that only support 32-bit instructions (as used in this class)
- RISC-V conditional branches can only reach $\pm 2^{10} \times 32$ -bit instructions on either side of PC



RISC-V B-Format for Branches



- B-format is mostly same as S-Format, with two register sources (rs1/rs2) and a 12-bit immediate imm[12:1]
- But now immediate represents values -4096 to +4094 in 2-byte increments
- The 12 immediate bits encode *even* 13-bit signed byte offsets (lowest bit of offset is always zero, so no need to store it)



Branch Example, Determine Offset

•RISC-V Code:

```
Loop: beq    x19, x10, End
      add    x18, x18, x10
      addi   x19, x19, -1
      j      Loop
End:   # target instruction
```

1
2
3
4

Count instructions from branch

- Branch offset = **$4 \times 32\text{-bit instructions} = 16\text{ bytes}$**
- (Branch with offset of 0, branches to itself)



Branch Example, Determine Offset

•RISC-V Code:

```
Loop: beq x19, x10, End
      add x18, x18, x10
      addi x19, x19, -1
      j Loop
```

```
End: # target instruction
```

1
2
3
4

Count
instructions
from branch

???????	01010	10011	000	?????	1100011
imm	rs2=10	rs1=19	BEQ	imm	BRANCH

Branch Example, Encode Offset

•RISC-V Code:

```
Loop: beq    x19, x10, End
      add    x18, x18, x10
      addi   x19, x19, -1
      j      Loop
End:  # target instruction
```

offset = 16 bytes = 8x2 bytes

???????	01010	10011	000	?????	1100011
imm	rs2=10	rs1=19	BEQ	imm	BRANCH

RISC-V Immediate Encoding

Instruction encodings, inst[31:0]

31	30	25	24	20	19	15	14	12	11	8	7	6	0														
funct7					rs2					rs1					funct3				rd			opcode				R-type	
imm[11:0]										rs1					funct3				rd			opcode				I-type	
imm[11:5]					rs2					rs1					funct3				imm[4:0]				opcode				S-type
imm[12 10:5]					rs2					rs1					funct3				imm[4:1 11]				opcode				B-type

Branch Example, complete encoding

beq **x19, x10,** offset = 16 bytes

13-bit immediate, imm[12:0], with value 16

0000000010000

imm[0] discarded,
always zero

imm[12]

imm[11]

0	000000	01010	10011	000	1000	0	1100011
---	--------	-------	-------	-----	------	---	---------

imm[10:5] rs2=10 rs1=19 BEQ imm[4:1] BRANCH

All RISC-V Branch Instructions

<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>000</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BEQ
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>001</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BNE
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>100</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BLT
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>101</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BGE
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>110</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BLTU
<code>imm[12 10:5]</code>	<code>rs2</code>	<code>rs1</code>	<code>111</code>	<code>imm[4:1 11]</code>	<code>1100011</code>	BGEU

Questions on PC-addressing

- **Does the value in branch immediate field change if we move the code?**
 - If moving individual lines of code, then yes
 - If moving all of code, then no ('position-independent code')
- **What do we do if destination is $> 2^{10}$ instructions away from branch?**
 - Other instructions save us

Questions on PC-addressing


- Does the value in branch immediate field change if we move the code?
 - If moving individual lines of code, then yes
 - If moving all of code, then no (because PC-relative offsets)
- What do we do if destination is $> 2^{10}$ instructions away from branch?
 - Other instructions save us

```
beq x10,x0,far  
# next instr
```

→

```
bne x10,x0,next  
j    far  
next: # next instr
```

U-Format for “Upper Immediate” Instructions



- Has 20-bit immediate in upper 20 bits of 32-bit instruction word
- One destination register, rd
- Used for two instructions
 - LUI – Load Upper Immediate
 - AUIPC – Add Upper Immediate to PC



LUI to Create Long Immediates

- LUI writes the upper 20 bits of the destination with the immediate value, and clears the lower 12 bits.
- Together with an ADDI to set low 12 bits, can create any 32-bit value in a register using two instructions (LUI/ADDI).

LUI x10, 0x87654 # x10 = 0x87654000

ADDI x10, x10, 0x321 # x10 = 0x87654321

One Corner Case

How to set 0xDEADBEEF?

LUI x10, 0xDEADB # x10 = 0xDEADB000

ADDI x10, x10, 0xEEF # x10 = 0xDEADAEEF

ADDI 12-bit immediate is always sign-extended, if top bit is set, will subtract -1 from upper 20 bits

Solution

How to set 0xDEADBEEF?

LUI x10, 0xDEADC # x10 = 0xDEADC000

ADDI x10, x10, 0xEEF # x10 = 0xDEADBEEF

Pre-increment value placed in upper 20 bits, if sign bit will be set on immediate in lower 12 bits.

Assembler pseudo-op handles all of this:

li x10, 0xDEADBEEF # Creates two instructions



Administrivia

- **Please don't come to labs asking 'please fix my code' or 'can we go over my code?'**
 - **Debug yourself!**
 - **Separate unit tests from functional/integration tests**
 - **If really needed, a question like "My code fails this test in this line, can you please help me find why?" has a better chance of getting the right result (working code)**

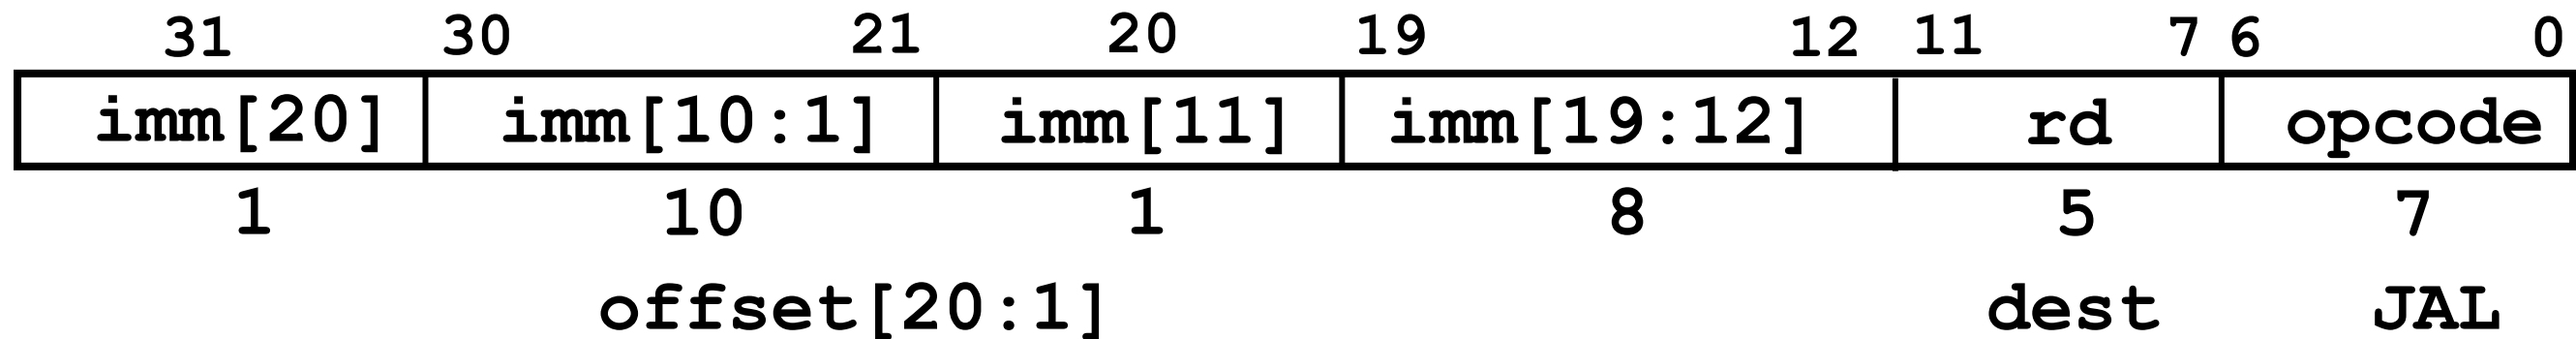


AUIPC

- Adds upper immediate value to PC and places result in destination register
- Used for PC-relative addressing

Label: AUIPC x10, 0 # Puts address of label in x10

J-Format for Jump Instructions



- **JAL saves PC+4 in register rd (the return address)**
 - Assembler “j” jump is pseudo-instruction, uses JAL but sets rd=x0 to discard return address
- **Set PC = PC + offset (PC-relative jump)**
- **Target somewhere within $\pm 2^{19}$ locations, 2 bytes apart**
 - $\pm 2^{18}$ 32-bit instructions
- **Immediate encoding optimized similarly to branch instruction to reduce hardware cost**

Uses of JAL

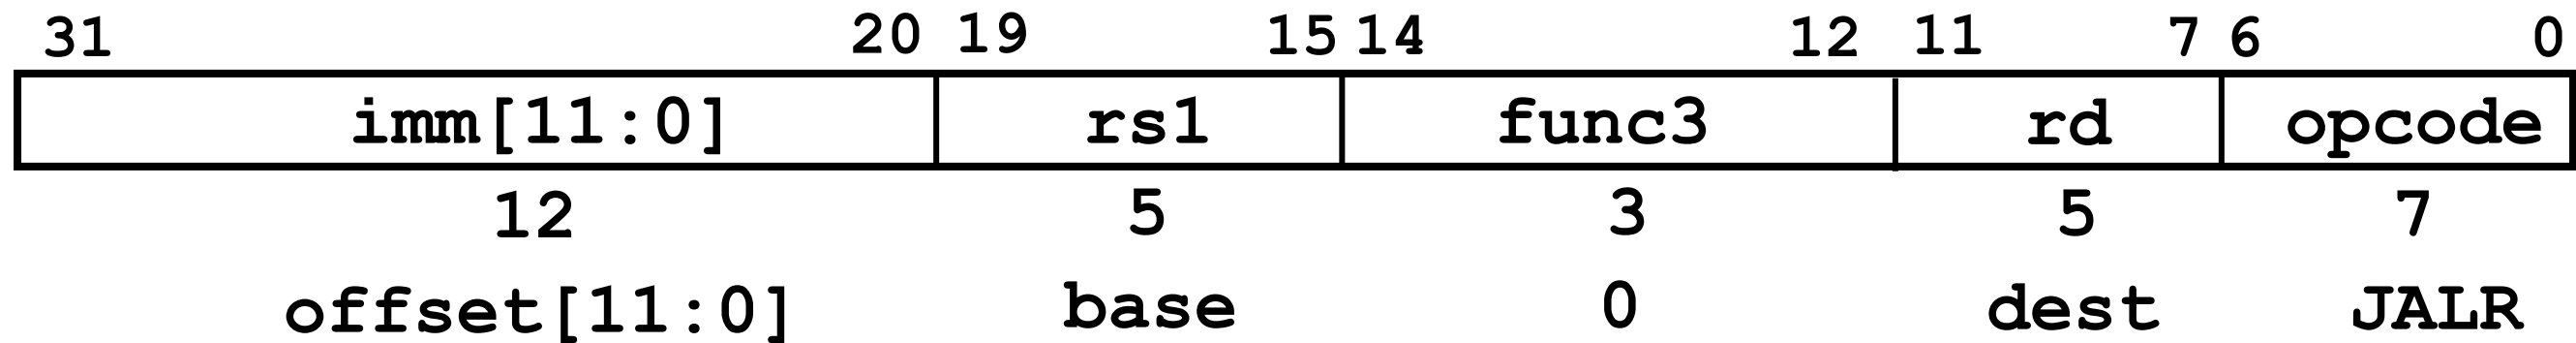
j pseudo-instruction

j Label = jal x0, Label # Discard return address

Call function within 2^{18} instructions of PC

jal ra, FuncName

JALR Instruction (I-Format)



- **JALR rd, rs, immediate**
 - Writes PC+4 to rd (return address)
 - Sets PC = rs + immediate
 - Uses same immediates as arithmetic and loads
 - **no** multiplication by 2 bytes
 - In contrast to branches and JAL

Uses of JALR

ret and jr psuedo-instructions

ret = jr ra = jalr x0, ra, 0

Call function at any 32-bit absolute address

lui x1, <hi20bits>

jalr ra, x1, <lo12bits>

Jump PC-relative with 32-bit offset

auipc x1, <hi20bits>

jalr x0, x1, <lo12bits>

Summary of RISC-V Instruction Formats

31	30	25	24	21	20	19	15	14	12	11	8	7	6	0			
funct7				rs2			rs1		funct3			rd			opcode		R-type
imm[11:0]						rs1		funct3			rd			opcode		I-type	
imm[11:5]				rs2			rs1		funct3			imm[4:0]			opcode		S-type
imm[12 10:5]				rs2			rs1		funct3			imm[4:1 11]			opcode		B-type
imm[31:12]									rd			opcode			U-type		
imm[20 10:1 11]]						imm[19:12]				rd			opcode			J-type	

Complete RV32I ISA

imm[31:12]				rd	0110111
imm[31:12]				rd	0010111
imm[20:10:11:19:12]				rd	1101111
imm[11:0]		rs1	000	rd	1100111
imm[12:10:5]	rs2	rs1	000	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	001	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	100	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	101	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	110	imm[4:1:11]	1100011
imm[12:10:5]	rs2	rs1	111	imm[4:1:11]	1100011
imm[11:0]		rs1	000	rd	0000011
imm[11:0]		rs1	001	rd	0000011
imm[11:0]		rs1	010	rd	0000011
imm[11:0]		rs1	100	rd	0000011
imm[11:0]		rs1	101	rd	0000011
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011
imm[11:0]		rs1	000	rd	0010011
imm[11:0]		rs1	010	rd	0010011
imm[11:0]		rs1	011	rd	0010011
imm[11:0]		rs1	100	rd	0010011
imm[11:0]		rs1	110	rd	0010011
imm[11:0]		rs1	111	rd	0010011

LUI
AUIPC
JAL
JALR
BEQ
BNE
BLT
BGE
BLTU
BGEU
LB
LH
LW
LBU
LHU
SB
SH
SW
ADDI
SLTI
SLTIU
XORI
ORI
ANDI

0000000		shamt	rs1	001	rd	0010011
0000000		shamt	rs1	101	rd	0010011
0100000		shamt	rs1	101	rd	0010011
0000000		rs2	rs1	000	rd	0110011
0100000		rs2	rs1	000	rd	0110011
0000000		rs2	rs1	001	rd	0110011
0000000		rs2	rs1	010	rd	0110011
0000000		rs2	rs1	011	rd	0110011
0000000		rs2	rs1	100	rd	0110011
0000000		rs2	rs1	101	rd	0110011
0100000		rs2	rs1	101	rd	0110011
0000000		rs2	rs1	110	rd	0110011
0000000		rs2	rs1	111	rd	0110011
0000	pred	succ	00000	000	00000	0001111
0000	0000	0000	00000	001	00000	0001111
0000000000000			00000	000	00000	1110011
0000000000001			00000	000	00000	1110011
csr			rs1	001	rd	1110011
csr			rs1	010	rd	1110011
csr			rs1	011	rd	1110011
csr			zimm	101	rd	1110011
csr			zimm	110	rd	1110011
csr			zimm	111	rd	1110011

SLLI
SRLI
SRAI
ADD
SUB
SLL
SLT
SLTU
XOR
SRL
SRA
OR
AND
FENCE
FENCE.I
ECALL
EBREAK
CSRRW
CSRWS
CSRRC
CSRRWI
CSRRSI
CSRRCI

Not in CS61C



“And in Conclusion...”

- **We have covered all RISC-V instructions and registers**
 - R-type, I-type, S-type, B-type, U-type and J-type instructions
 - Practice assembling and disassembling

